

OPS-WARM-OUTREACH

Master Context Prompt & Build Plan

Assigned to	Kyle (Lead Developer)
Direction	Jon (Founder, FascinateCopy)
Ops	Julia (Director of Operations)
Repo (proposed)	Noah-dev2217/ops-warm-outreach
Infrastructure	Hostinger VPS · n8n · Postgres · Docker · Anthropic API
Objective	95% automated client acquisition for \$5K–\$7K/mo retainers

How to use this document: Kyle — paste this entire PDF's contents (or attach the file) into a fresh Claude conversation as your first message. Claude will have full context on the methodology, technical stack, phase order, and resolved decisions. You can then ask for specific tickets, Python code, n8n JSON, or Claude prompts for any phase and Claude will answer without needing re-briefing.

Pattern reference: This build follows the same conventions as OPS-61 and FC IG Lead Discovery — per-phase build specs, HITL gates, no automated closes, Postgres-backed state, n8n orchestration. If you've touched those repos, this will feel familiar.

System Role & Objective (for Claude)

You are a Lead AI Automation Engineer, a Technical Product Manager, and a High-Ticket Growth Strategist. Your role is to help me (Kyle) architect, plan, and build a 95% automated multi-channel client acquisition system for FascinateCopy.

This system targets high-net-worth B2B founders — info businesses, coaches, course creators, and qualifying ecommerce brands doing \$100K+/month — to sell them \$5,000–\$7,000/month Full Stack Email Marketing retainers.

You must adhere strictly to the methodology, the psychological copywriting frameworks, the operational rules, the technical stack, and the 8-phase pipeline below. Once you have internalized this document, help me write developer-ready tickets, n8n workflow architectures, Python scripts, and Claude API prompts to build this system phase by phase.

PART 1 — Core Methodology (The Hybrid System)

1A. Operational Spine (Agency Model)

ICP — Strict Dual-Tier Filtering

Primary Target (Info / Coaches / Courses): \$100K+/month gross revenue with a 10K+ engaged email list. Digital products and high-ticket coaching (\$2K–\$5K+ programs) run 60–80%+ profit margins, so these founders can absorb \$5–7K retainers easily. This is the goldmine.

Secondary Target (Ecom, gated tighter): 80K+ list with \$70+ AOV, OR 60K–100K list with \$27+ AOV. Ecom margins (20–40%) mean brands often cap retainers at \$1–3K, so thresholds must be higher to justify \$5K+.

Hard Disqualifiers (auto-archive): dropshippers, brick-and-mortar businesses, unproven-revenue startups, any lead scoring under 10/25 on the qualification rubric.

25-Point Qualification Scoring

Prospects graded 1–5 on five metrics: Revenue level, Profit margin, Email list size, Growth rate, Decision-making authority. Routing: **20+** = dream client, high-priority. **15–19** = queued warm. **10–14** = queued cold, low priority. **<10** = auto-archived.

Core Value Prop — The 17-Step Email Profit Formula

We do not just write copy — we build complete revenue systems: 30 daily broadcast emails per month + 10 automated workflow emails (welcome sequences, abandoned cart, browse abandonment, upsells / downsells) + deliverability optimization + weekly quantitative and qualitative reporting.

The Hook — The Automated Audit

We run automated Technical Funnel Audits. We opt into the prospect's funnels, abandon a high-intent action, map their automated emails over 72 hours, and identify exact revenue gaps where they are bleeding money. This audit becomes the personalization payload for every DM.

The Close Mechanism

We **NEVER** close in DMs. The entire system's goal is to push prospects to a 20-minute video screen-share call, where we present the funnel audit visually, pitch the \$5K+ retainer, and use a dead-silence closing technique.

1B. Communication Supplement (DM Copy Layer)

- **160-Character Rule:** All initial outreach DMs must be under 160 characters. No blocky agency-spam paragraphs.
- **S.I.P.E. Framework:** Initial hooks must be **Short**, **Incomplete**, **Personal**, **Emotional**.
- **Tease-Then-Deliver:** Do not dump audit findings in message 1. Tease the specific gap, ask permission, deliver the value only after they reply yes.
- **Ladder of Tiny Yeses:** Escalate through micro-commitments instead of asking for a 45-min call upfront. E.g. "want me to send a 2-min video breakdown?" → "would it help if we looked at this together on a quick screen share?"

1C. Strict Contradictions to Avoid

You must completely **DISCARD** the following common marketing tropes — they conflict with selling premium \$5K+ retainers to sophisticated founders:

- **NO closing in DMs.** Escalate to a video screen-share.
- **NO trailing-off cliffhangers.** Premium founders hate curiosity bait. We tease the audit, but deliver the exact technical diagnosis immediately once they reply.
- **NO fake AI personalization.** Do not generate compliments about their content ("love your recent post about leadership!"). Personalization must come from actual technical data gathered during the audit.

PART 2 — Technical Stack

Already Have (reuse directly)

Infrastructure	Hostinger VPS · n8n (self-hosted) · Postgres · Supabase · Docker
AI	Anthropic API — Haiku (classify/score) · Sonnet (draft) · Opus (high-value drafts)
Cold Email Stack	15 inboxes across fascinatecopy.site · fascinatehq.online · usefascinate.space
HITL Pattern	Discord bot with Approve/Edit/Reject buttons (from FC IG Lead Discovery / OPS-61)
Browser Automation	In-house Playwright experience (Loomless, FC IG Auto-DM, OPS-52)

Build In-House

- **Audit worker** — Python 3.11 + Playwright + patchright stealth + imap-tools (the moat)
- **DM sending layer** — Playwright routines through Dolphin profiles + IPRoyal proxies (in-house, no PhantomBuster)
- **Discord HITL bot** — reuse FC IG Lead Discovery pattern
- **Scoring engine** — n8n + Claude Haiku applying 25-point rubric
- **DM drafting engine** — n8n + Claude Sonnet / Opus with S.I.P.E. prompt library
- **Reply triage classifier** — n8n + Haiku (interested / objection / not_interested / unclear)
- **Postgres schema** — leads, lead_enrichment, lead_audits, lead_outreach, outreach_events
- **Small scrapers** — Meta Ad Library, social posts, Similarweb free tier

Buy (Cheapest Viable Tier)

Component	Vendor	Cost
Residential proxies	IPRoyal pay-as-you-go (\$7/GB)	~\$20–40/mo
Anti-detect browser	Dolphin{anty} free tier (10 profiles)	\$0
B2B contact data	Apollo free tier → \$49/mo Basic when needed	\$0 → \$49
Scraper actors (backup)	Apify pay-as-you-go	~\$5–15/mo
Analytics	Metabase OSS self-hosted	\$0
Call booking	Cal.com + Google Meet	\$0

Invest Time (no money, no shortcut)

- **Warm social accounts yourself** — 3 IG, 2 LinkedIn, 2 FB personas. 30–45 days of human-looking activity before any DM sends. Start Day 1, parallel to build.
- **Self-created ProtonMail inbox pool** — 20+ disposable inboxes with IMAP enabled for the audit worker.

Starting monthly cost: ~\$25–55. Everything else is developer time on infrastructure already owned.

Skip Entirely

- **SmartLead / Instantly** — the existing 15-inbox n8n cold email stack replaces these.
- **PhantomBuster / Apify DM actors** — build DM sending in-house on Playwright.
- **Streamlit** — Metabase is faster for MVP.
- **Clay** — pricier than Apollo, not needed at this stage.

PART 3 — Build Order (Phases 0–8)

Each phase is a self-contained ticket in the format Kyle knows: **Task Title · Task Description · Expected Outcome**. Build in order. Do not skip Phase 0 — every downstream phase depends on it.

■ PHASE 0 — Foundation & Bottleneck Unblock

Timeline: Week 1

Task Title

Foundation & Bottleneck Unblock

Task Description

Everything downstream depends on this phase. Two clocks start Day 1 — social account warmup (30–45 days) and cold email domain warmup (14 days, ongoing on existing 15-inbox stack). Neither has a shortcut. Provision the following in parallel:

1. **Postgres schema** on existing DB — create tables: leads, lead_enrichment, lead_audits, lead_outreach, outreach_events. Include audit-state fields (audit_status: pending / in-progress / complete / failed) — 72hr latency is a real state machine, not a synchronous call.
2. **IPRoyal account** — pay-as-you-go residential proxies, credentials stored in n8n vault.
3. **Dolphin{anty} free tier** — install on ops machine, provision one browser profile per planned social account.
4. **Social account warmup** — 3 IG, 2 LinkedIn, 2 FB personas. Human-looking activity: follow, like, comment, post, DM real people. 30–45 day timer starts now. Reuse persona resolutions from prior warm-DM Phase 0 planning.
5. **ProtonMail inbox pool** — 20 accounts minimum, self-created, credentials in vault, IMAP enabled and verified.
6. **Metabase container** — Docker on VPS, points at Postgres, reachable behind Traefik or basic auth.
7. **Cold email stack confirmation** — verify existing 15-inbox warmup is healthy, reserve capacity for parallel sending.

Expected Outcome

Schema live. Proxies + Dolphin operational. Metabase reachable. Warmup clocks ticking on social accounts and inboxes. All infrastructure ready to accept traffic from Phase 3 onward.

■ PHASE 1 — Sourcing Pipeline

Timeline: Week 1–2

Task Title

Sourcing Pipeline

Task Description

Build n8n workflow that pulls prospects matching the strict dual-tier ICP into the Leads table. Sources:

- **Apollo free tier** — founder contacts, company revenue signal, social handles (IG, LinkedIn).
- **Custom Playwright scraper (in-house)** — Meta Ad Library, confirms prospect is actively spending on ads.

- **StoreLeads free tier / manual export** — ecom side only, filter for 80K+ list + \$70+ AOV or 60K–100K + \$27+ AOV.

- **Info-business sourcing** — creator databases, coach directories, course marketplaces. Build a small ingest script per source.

Workflow must apply hard disqualifiers at ingestion (dropshippers, brick-and-mortar, unproven revenue) before writing to Leads. Deduplicate by domain + founder email.

Expected Outcome

leads table populated with 100–300 pre-filtered prospects per day. Zero dropshippers. Zero local businesses. Zero sub-threshold lists.

■ PHASE 2 — Enrichment Pipeline

Timeline: Week 2

Task Title

Enrichment Pipeline

Task Description

n8n workflow triggered by new Leads rows. For each lead:

- **Similarweb free tier scrape** (custom Playwright) — traffic estimate.
- **BuiltWith free data / paid API at scale** — tech stack (Klaviyo, ConvertKit, Shopify Plus, Kajabi, ClickFunnels, etc.).
- **Custom Playwright social post scraper** — 3 most recent posts from prospect's IG + LinkedIn, stored as raw text for later DM personalization.
- **Funnel entry point detection** — locate lead magnet, opt-in form, application form, or product page. This URL feeds the audit worker.

Write everything to `lead_enrichment`. Flag leads missing critical enrichment data (e.g. no discoverable funnel) as `audit_status = failed` with reason. Do not silently drop.

Expected Outcome

Every viable lead has a complete enrichment row with tech stack, traffic estimate, funnel entry URL, and 3 recent social posts. Failed enrichments logged with reason.

■■ PHASE 3 — Python Audit Worker — THE MOAT

Timeline: Week 2–4

Task Title

Python Audit Worker — THE MOAT

Task Description

This is the differentiator. Dockerized Python 3.11 service running on VPS, orchestrated by n8n triggers.

Two audit modalities:

- **Ecom mode** — Playwright + patchright stealth: navigate to product page, add to cart, begin checkout with disposable ProtonMail address, abandon. IPRoyal proxy per session.

- **Info-business mode** — Playwright: opt into main lead magnet OR start "book a call" / application flow, submit email, exit before completion.

Wait & parse: Worker sets `audit_status = in-progress` with `audit_completes_at = now() + 72hr`. n8n scheduler picks up ready-to-parse rows every hour; worker connects via IMAP (imap-tools) to the disposable inbox and parses received mail.

Structured JSON output to lead_audits: `welcome_email_received` (bool), `welcome_has_cta` (bool), `abandoned_cart_count` (int), `abandoned_application_followup_count` (int), `first_recovery_delay_hours` (float), `discount_offered` (bool), `discount_amount` (string|null), `total_emails_received_72h` (int), `deliverability_signal` (primary / promotions / spam), `gap_summary` (Claude-generated one-liner).

Critical: treat "audit-pending" as a first-class state. Phase 4 scoring must not run on incomplete audits. Failed audits (site down, no email received, checkout broken) must log a reason, not silently succeed with zeros.

Expected Outcome

Worker produces structured audit JSON for both ecom and info-business funnels within 72–80hr of trigger. Failure modes logged with cause. Rotating inbox + proxy pool prevents cross-contamination between audits.

■ PHASE 4 — Qualification & Scoring

Timeline: Week 4

Task Title

Qualification & Scoring

Task Description

n8n workflow triggered when `lead_audits.audit_status = complete`. Pulls enrichment + audit rows and sends a structured payload to Claude Haiku with the 5-factor 25-point rubric (Revenue, Profit margin, List size, Growth rate, Decision authority).

- **<10** → auto-archive with reason.
- **10–14** → queued cold, low priority.
- **15–19** → queued warm.
- **20+** → high-priority flag, immediately routes to Phase 5.

Score, sub-scores, and Claude's reasoning written to `leads.qualification_score` and `leads.qualification_reasoning`.

Expected Outcome

Every audited lead has a numeric score and archival/routing decision within seconds of audit completion. Reasoning auditable in the DB, not opaque.

■ PHASE 5 — AI Personalization & DM Drafting

Timeline: Week 4–5

Task Title

AI Personalization & DM Drafting

Task Description

n8n workflow triggered by score ≥ 15 . Sends enrichment + audit + recent social posts to Claude Sonnet (score 15–19) or Opus (score 20+).

Prompt library enforces:

- **Message 1:** ≤ 160 chars, S.I.P.E. framework, references specific audit gap, references prospect's recent social post (not a generic compliment).
- **Message 2:** tease-then-deliver — delivers the exact gap technically.
- **Message 3:** soft escalation to screen-share via Ladder of Tiny Yeses.
- No cliffhangers, no fake compliments, no "close" language.

Parallel cold email variant: Same audit-driven personalization, 4-touch sequence, plain-text style, throttled to existing 15-inbox stack (20/day ramping to 50/day per inbox).

All drafts written to `Lead_outreach` with `channel` (`dm_ig` / `dm_linkedin` / `dm_fb` / `cold_email`) and `status` = `draft`.

Expected Outcome

Every qualified lead has channel-appropriate drafts stored in DB, ready for HITL review. DMs pass character and framework validation before saving.

■■ PHASE 6 — Discord HITL Bot

Timeline: Week 5

Task Title

Discord HITL Bot

Task Description

Reuse the working HITL pattern from FC IG Lead Discovery (M8 module) and OPS-61. Bot posts each draft to a Discord channel with:

- Lead summary (company, score, gap headline).
- Full drafted message.
- Buttons: **[Approve]** **[Edit]** **[Reject]** **[Skip Lead]**.
- Edit opens a modal for inline rewrite.

Approved messages flip `Lead_outreach.status` = `approved` and queue for Phase 7. Rejected drafts log a reason for prompt improvement. High-value leads (score 20+) require explicit approval; score 15–19 can be configured for optional bulk-approve mode.

Expected Outcome

Julia can review and dispatch 50–100 drafts/day in under 30 minutes. No message sends without explicit approval on high-value leads.

■ PHASE 7 — Sending & Reply Triage

Timeline: Week 5–6

Task Title

Sending & Reply Triage

Task Description

DM sending (in-house Playwright, no PhantomBuster): Playwright routines drive Dolphin browser profiles through IPRoyal proxies. **Hard cap: 10–20 DMs/day/profile.** Human-timing jitter, session rotation, cooldown on soft-block signals (reuse pattern from FC IG Lead Discovery M5 executor).

Cold email sending: Dispatch through existing 15-inbox n8n stack. 20/day/inbox ramp, cap 50/day/inbox. Uses same audit-derived personalization as DMs.

Reply triage: Webhook or IMAP poll triggers n8n on inbound. Claude Haiku classifies as interested / objection / not_interested / unclear. Interested replies post to Discord with full context and ping human operator for takeover. Human executes Ladder of Tiny Yeses and books screen-share.

All sends and replies logged to outreach_events.

Expected Outcome

System sends 30–100 DMs/day + 300–500 cold emails/day within compliance limits. Interested replies land in Discord with sub-5-minute latency. Zero automated closes.

■ PHASE 8 — Analytics Dashboard

Timeline: Week 6

Task Title

Analytics Dashboard

Task Description

Metabase views over Postgres. Core metrics:

- Leads sourced / day, / channel
- Audits complete / failed / pending
- Score distribution
- DMs sent, replies, positive-reply rate
- Cold emails sent, replies, positive-reply rate
- Calls booked, cost per booked call
- Per-channel funnel conversion

Daily Discord digest posts key numbers at 8am. Weekly deep-dive report for Jon on Sundays.

Expected Outcome

Jon and Julia can answer "did we make progress this week" in under 30 seconds. Bottlenecks (audit failure rate, score distribution collapse, reply rate drop) surface visibly.

PART 4 — Critical Path Bottlenecks

Run these in parallel with Phases 0–2. If they lag, launch stalls regardless of code progress.

Bottleneck	Detail
Social account warmup	30–45 day human-looking activity clock. No shortcut. Start Day 1.
Cold email domain warmup	Ongoing on existing 15-inbox stack. Verify healthy before Phase 5 drafts hit it.
ProtonMail inbox pool	Must exist before Phase 3 runs on real leads. 20 accounts minimum.
Proxy + Dolphin setup	Must be live before any sending or audit traffic. Free-tier proxies = burned accounts.

PART 5 — Resolved Decisions

These were open at the start of planning and are now locked. Do not re-litigate; if a constraint conflicts with implementation, escalate to Jon rather than assume.

Decision	Resolution
ICP scope	Dual-tier. Primary = info businesses / coaches / courses at \$100K+/mo + 10K+ list. Secondary = ecom, gated tig
Audit modality	Dual mode. Ecom = cart-abandon flow. Info-business = opt-in / application-abandon flow. Same 72hr IMAP capt
Cold email role	Parallel channel using the existing 15-inbox n8n stack. Same audit-driven personalization as DMs. Throttled 20/
DM execution	In-house Playwright routines through Dolphin profiles + IPRoyal proxies. No PhantomBuster, no Apify DM actors
Assignee	Kyle leads. Rinoah is loaded on OPS-61 and not available for this build. Kevin available for support on pipeline v

Handoff Note — Kyle

You have full context. Recommended first move: open a fresh Claude conversation, paste this PDF as your first message, then ask Claude for:

- BUILD_SPEC_0_FOUNDATION.md** — Phase 0 ticket with acceptance criteria, files touched, LOE.
- Postgres schema SQL** — full DDL for all 5 tables with indexes and state fields.
- n8n workflow JSON scaffolds** for Phases 1 and 2 (sourcing + enrichment).
- Python audit worker scaffold** — Dockerfile, requirements.txt, entrypoint, and modality router.
- Claude prompt library** — S.I.P.E. draft prompt, scoring rubric prompt, reply-triage prompt.

Conventions to follow (matches OPS-61 and FC IG Lead Discovery):

- No secrets in the repo. Everything sensitive in .env, gitignored.
- Per-phase README.md with run instructions and acceptance criteria.
- Per-phase requirements.txt or package.json — isolate Python services.
- n8n workflows exported as JSON to shared/n8n-templates/ for version control.

- Postgres migrations in `db/migrations/`, numbered sequentially.
- One Claude Code session per phase — never one-shot the whole system.

If you get stuck: escalate to Jon in Discord with the specific decision needed. Do not assume scope; do not silently expand. Jon prefers direct questions over speculative builds.